



Time-series,  
Spring, 2026



# Python & Time Series Fundamentals

*Faculty of DS & AI  
Spring semester, 2026*

Trong-Nghia Nguyen

# Content

- NumPy Review for Time Series
- Pandas for Time Series
- Technical Indicators and EDA

# Content

- NumPy Review for Time Series
- Pandas for Time Series
- Technical Indicators and EDA

# NumPy Review for Time Series

## NumPy arrays

### Exercise:

1. Create a NumPy array with 50 evenly spaced values from 0 to 10 (inclusive). What is the shape and dtype?
2. Given the array `arr = np.array([10, 20, 30, 40, 50, 60, 70, 80, 90, 100])`:
  - Extract the first 5 elements
  - Extract elements from index 3 to 7 (inclusive)
  - Extract every other element starting from index 0
3. Create two arrays: `a = np.array([1, 2, 3])` and `b = np.array([10, 20, 30])`.
  - Add them elementwise
  - Multiply them elementwise
  - What happens if you try `a + 5`? (Broadcasting)
4. Create a time series array with 100 daily values starting from January 1, 2020. Use `np.arange()` to create day numbers (0 to 99), then create values as `100 + 2day + np.random.normal(0, 5, 100)`. What is the mean and standard deviation?

# NumPy Review for Time Series

## NumPy Operations for Time Series

General representation of time series data:

$$\mathbf{X} = \{x_t \in \mathbb{R}^d \mid t=1, 2, \dots, T\} \quad (1)$$

- $\mathbf{X}$ : the complete time-series dataset
- $t$ : time index
- $T$ : total number of observed time steps
- $x_t$ : feature vector at time step  $t$
- $d$ : dimensionality of the feature space
- $\mathbb{R}^d$ :  $d$ -dimensional real-valued space

**Interpretation:**

The time series is modeled as a sequence of  $d$ -dimensional observations indexed by time.

# NumPy Review for Time Series

## NumPy Operations for Time Series

Time-series decomposition:

$$x_t = \tau_t + s_t + c_t + \epsilon_t \quad (2)$$

- $x_t$ : observed value at time step  $t$
- $\tau_t$  (**trend**): long-term progression of the series
- $s_t$  (**seasonality**): repeating patterns with a fixed period
- $c_t$  (**cyclic component**): non-periodic, long-term oscillations
- $\epsilon_t$  (**noise**): random fluctuations or unexplained residuals

**Interpretation:**

Each observation is decomposed into interpretable structural components plus random noise.

# NumPy Review for Time Series

## NumPy Operations for Time Series

Forecasting formulation:

$$\hat{x}_{t+1} = f(x_t, x_{t-1}, \dots, x_{t-k}) \quad (3)$$

- $\hat{x}_{t+1}$ : predicted value at the next time step
- $f(\cdot)$ : forecasting function (e.g., ARIMA, LSTM, Transformer)
- $k$ : look-back window size
- $x_{t-i}$ : historical observations used for prediction

**Interpretation:**

Future values are estimated based on a finite window of past observations.

# NumPy Review for Time Series

## NumPy Operations for Time Series

Multivariate time series:

$$x_t = \begin{bmatrix} v_t^{(1)} \\ v_t^{(2)} \\ \vdots \\ v_t^{(m)} \end{bmatrix}, X \in \mathbb{R}^{T \times m} \quad (4)$$

- $x_t$ : observation vector at time step  $t$
- $v_t^{(i)}$ : the  $i$ -th variable at time step  $t$
- $m$ : number of variables (features or channels)
- $T$ : number of time steps
- $X \in \mathbb{R}^{T \times m}$ : time–feature matrix representation of the dataset

Interpretation:

The time series is represented as a matrix, where rows correspond to time steps and columns correspond to variables.

# NumPy Review for Time Series

## NumPy Operations for Time Series

Scrolling (Sliding) Window Idea:

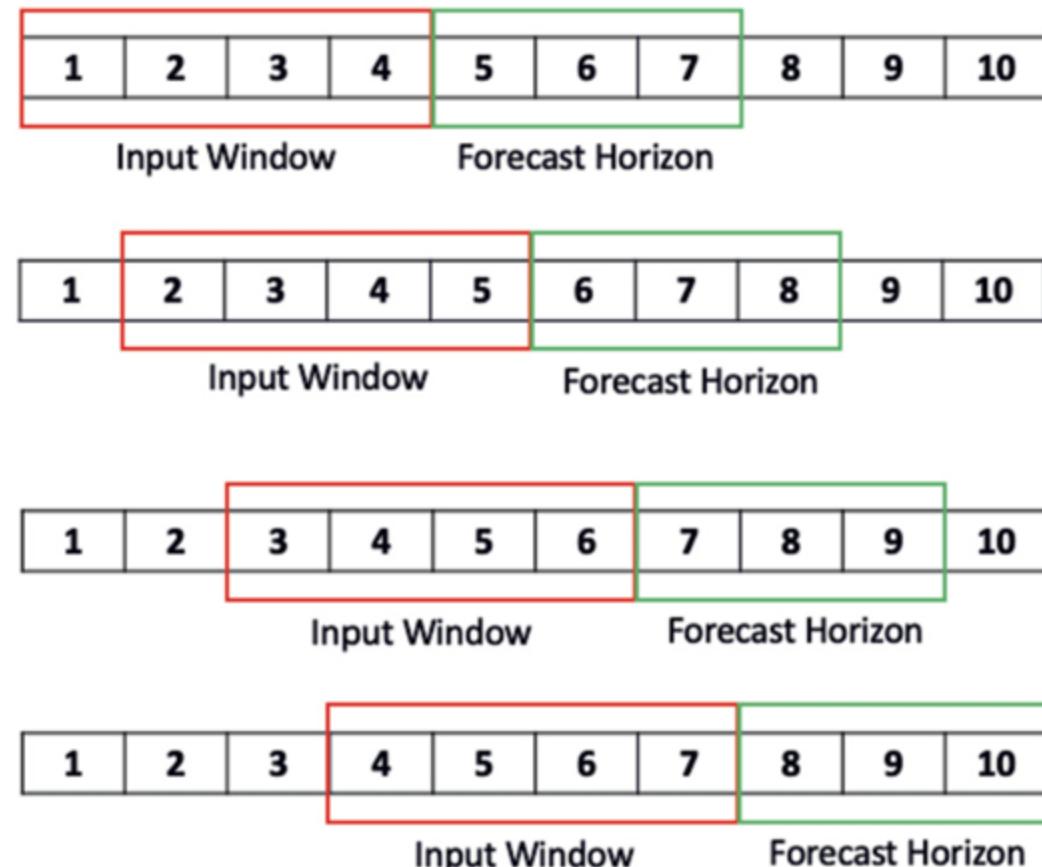
Choose a window size  $k$  (e.g.,  $k = 10$ )

Slide the window one step at a time along the series

At each position  $t$ , compute  $\bar{x}_t$  using the formula (5)

This produces a new, shorter series of smoothed values

```
# Rolling window calculation (manual)
window_size = 10
Kernel = np.ones(window_size)/window_size,
rolling_mean = np.convolve(values, kernel, mode = .
Valid)
```



# NumPy Review for Time Series

## NumPy Operations for Time Series

**Simple Moving Average (SMA)** over a window of size  $k$ :

$$\bar{x}_t = \frac{1}{k} \sum_{i=0}^{k-1} x_{t-i} \quad (5)$$

$k$  = window size

Larger  $k \rightarrow$  smoother curve

This averages the most recent  $k$  points  $x_t, x_{t-1}, \dots, x_{t-k+1}$  to smooth shortterm fluctuations.

# Content

- NumPy Review for Time Series
- Pandas for Time Series
- Technical Indicators and EDA

# Pandas for Time Series

## Pandas for Time Series

Dataset Description:

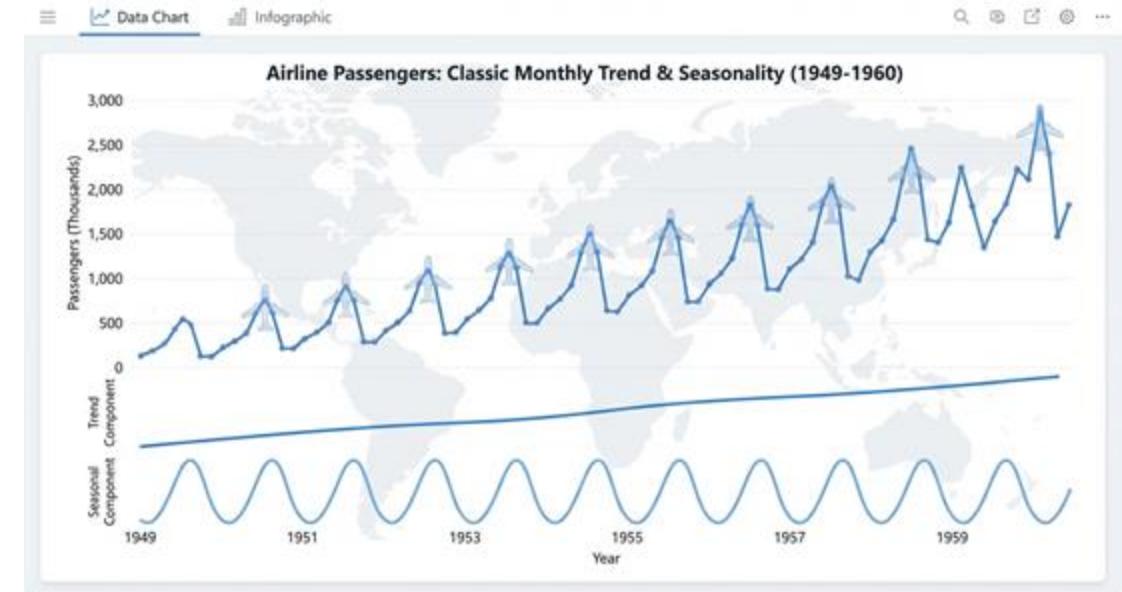
- Monthly airline passenger numbers (1949–1960)
- Univariate time series
- Clear trend and seasonality
- Widely used benchmark dataset in Time-series analysis

```
from statsmodels.datasets import get_rdataset
import pandas as pd

# Load AirPassengers dataset
data = get_rdataset('AirPassengers', 'datasets')
df = data.data

# Convert to Pandas Series with datetime index
ts = pd.Series(df['value'].values,
               index=pd.date_range('1949-01', periods=len(df), freq='M'))

print(ts.head())
print(ts.info())
```



AirPassengers dataset

# Pandas for Time Series

## Pandas for Time Series

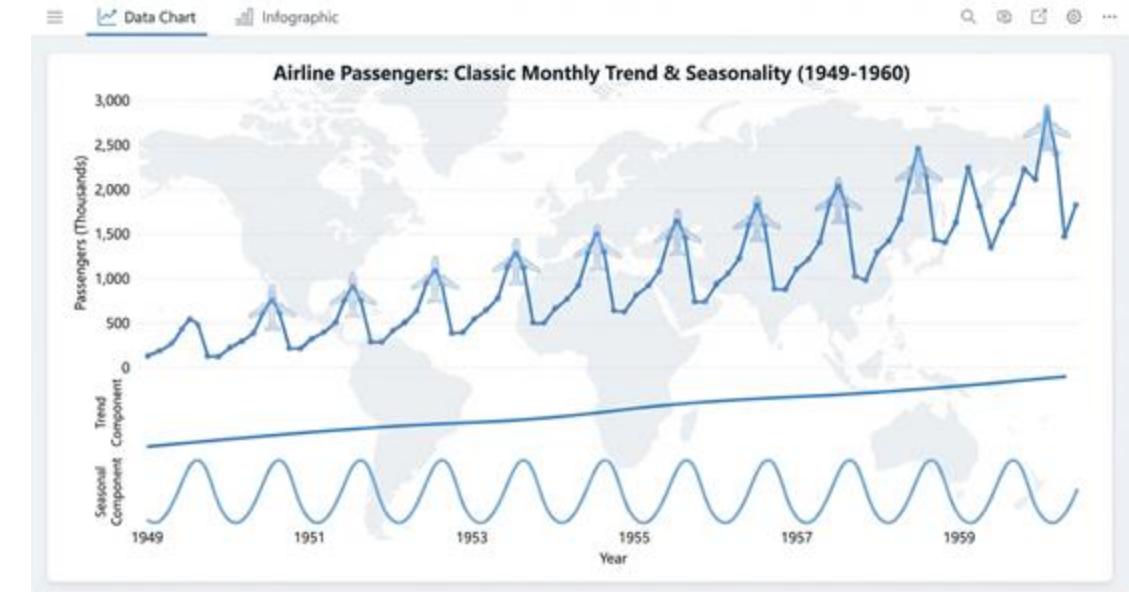
Time Indexing and Slicing:

- Datebased indexing
- Subsetting by year or date range
- Shifting the time series
- Differencing for trend removal

```
# Datebased indexing  
ts_1950 = ts['1950'] # All data from 1950  
ts_1950_1955 = ts['1950':'1955'] # Date range
```

```
# Shifting  
ts_lag1 = ts.shift(1) # Shift by 1 period
```

```
# Differencing (removes trend)  
ts_diff = ts.diff() # First difference
```



AirPassengers dataset

# Pandas for Time Series

## Pandas for Time Series

- Resampling:
  - Change frequency of observations
  - Aggregate data for analysis
  - Align multiple time series
  - Types:
    - Downsampling: Reduce frequency (e.g., daily → monthly)
    - Upsampling: Increase frequency (e.g., monthly → daily)
    - Common Frequencies:  
`'D' (daily), 'W' (weekly), 'M' (monthly), 'Q' (quarterly), 'Y' (yearly)`
    - Aggregation Functions:  
``mean()`, `sum()`, `last()`, `first()`, `max()`, `min()``

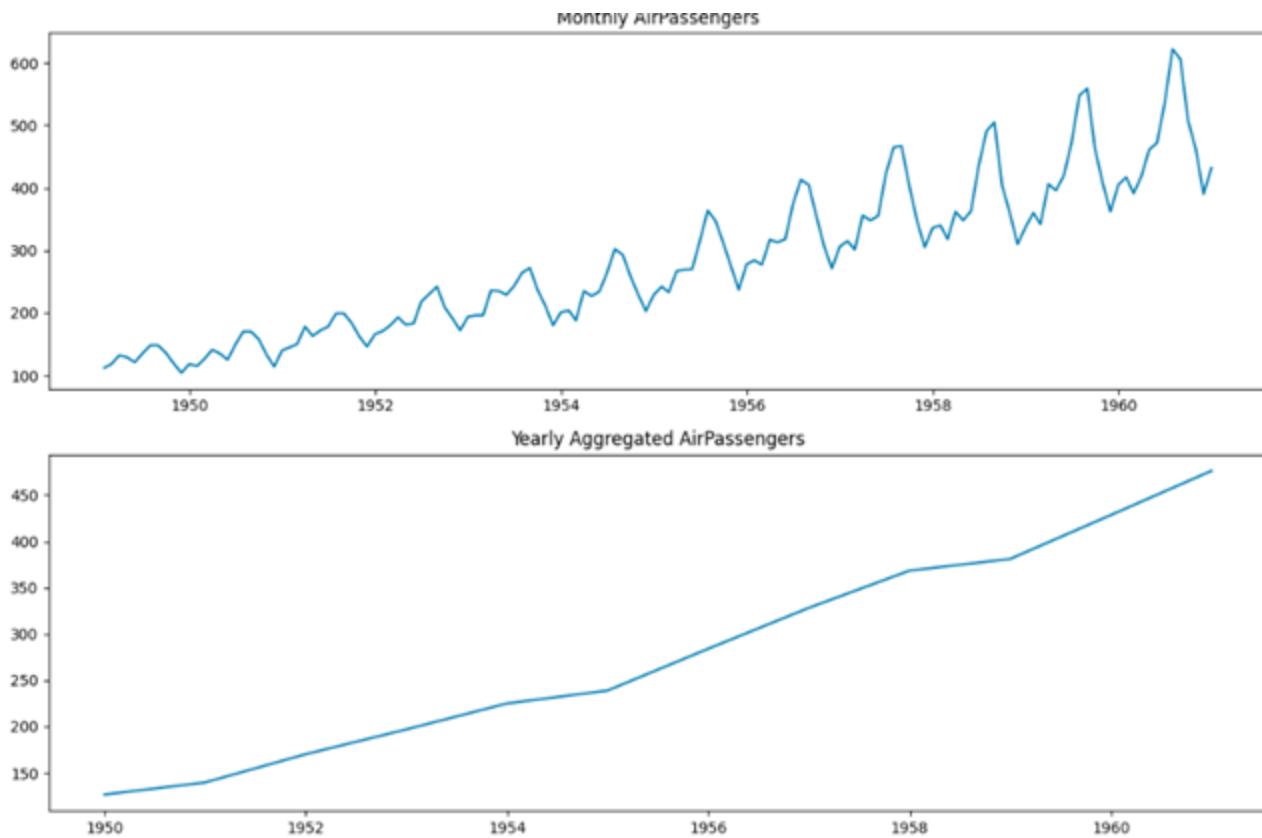
# Pandas for Time Series

## Pandas for Time Series

- Resampling AirPassengers
  - Monthly to yearly aggregation
  - Effect of resampling on trend and variance
  - Visual comparison of different frequencies

```
# Monthly to yearly aggregation  
  
yearly = ts.resample('Y').mean() # Yearly mean  
  
yearly_sum = ts.resample('Y').sum() # Yearly sum
```

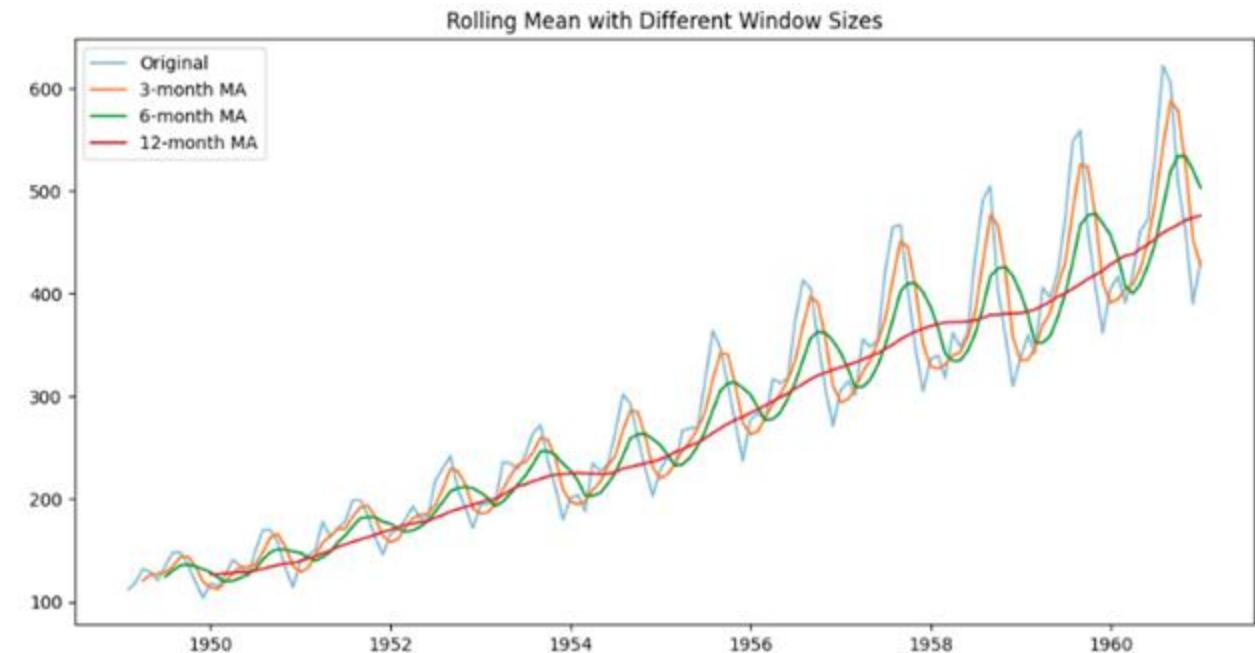
```
# Visual comparison  
  
import matplotlib.pyplot as plt  
  
fig, axes = plt.subplots(2, 1, figsize=(12, 8))  
  
axes[0].plot(ts.index, ts.values, label='Monthly')  
  
axes[0].set_title('Monthly AirPassengers')  
  
axes[1].plot(yearly.index, yearly.values, label='Yearly Mean')  
  
axes[1].set_title('Yearly Aggregated AirPassengers')  
  
plt.tight_layout()  
  
plt.show()
```



# Pandas for Time Series

## Pandas for Time Series

- Rolling Statistics with Pandas
  - Rolling mean
  - Rolling standard deviation
  - Effect of window size (3, 6, 12)
  - Connection to NumPy sliding window



# Pandas for Time Series

## Pandas for Time Series

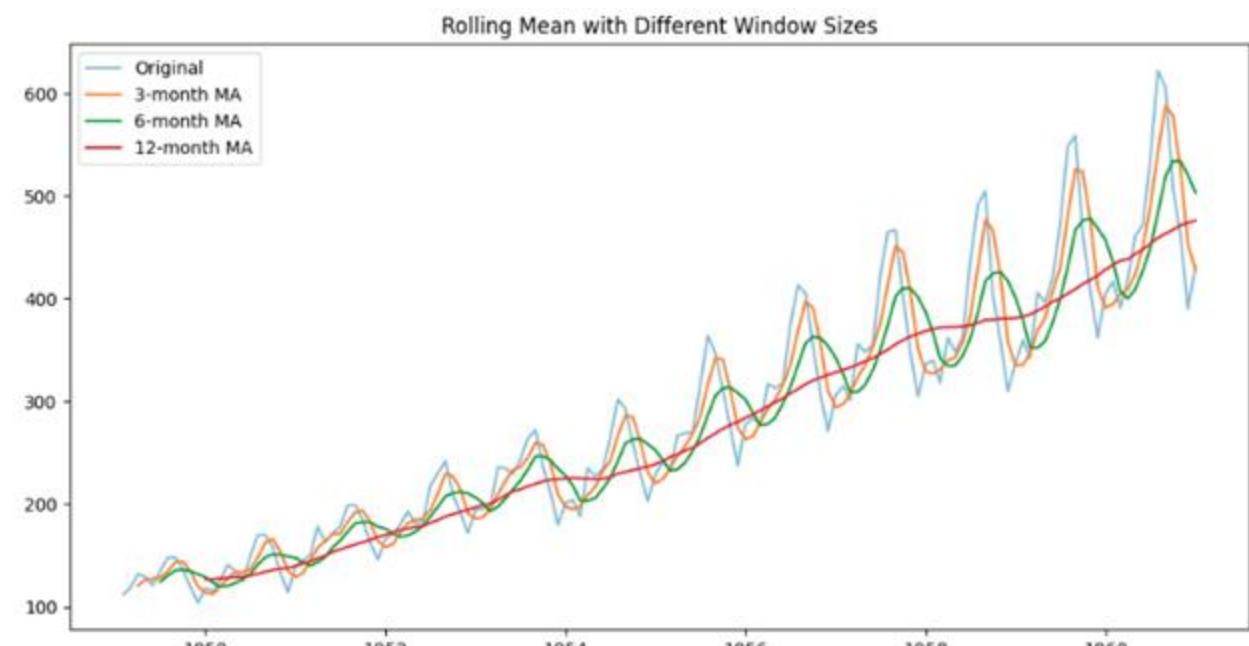
- Rolling Statistics with Pandas

```
# Rolling statistics
ts_rolling_mean_12 = ts.rolling(window=12).mean() # 12month rolling
ts_rolling_std_12 = ts.rolling(window=12).std() # 12month rolling

# Different window sizes
ts_rolling_3 = ts.rolling(window=3).mean()
ts_rolling_6 = ts.rolling(window=6).mean()
ts_rolling_12 = ts.rolling(window=12).mean()

# Visualization
plt.figure(figsize=(12, 6))

plt.plot(ts.index, ts.values, label='Original', alpha=0.5)
plt.plot(ts_rolling_3.index, ts_rolling_3.values, label='3month MA')
plt.plot(ts_rolling_6.index, ts_rolling_6.values, label='6month MA')
plt.plot(ts_rolling_12.index, ts_rolling_12.values, label='12month MA')
plt.legend()
plt.title('Rolling Mean with Different Window Sizes')
plt.show()
```



# Content

- NumPy Review for Time Series
- Pandas for Time Series
- Technical Indicators and EDA

# Technical Indicators and EDA

## From SMA to EMA

- Simple Moving Average (SMA): Equal weights for all observations in window.

$$\text{SMA}_n = \frac{1}{n} \sum_{i=1}^n P_i \quad (6)$$

- Exponential Moving Average (EMA): More weight to recent observations

$$\text{EMA}_t = \alpha P_t + (1 - \alpha) \text{EMA}_{t-1} \quad (7)$$

Where  $\alpha = 2 / (n + 1)$

EMA is more responsive to recent changes

SMA is smoother but less reactive

Smoothing any time series data

Trend detection in various domains

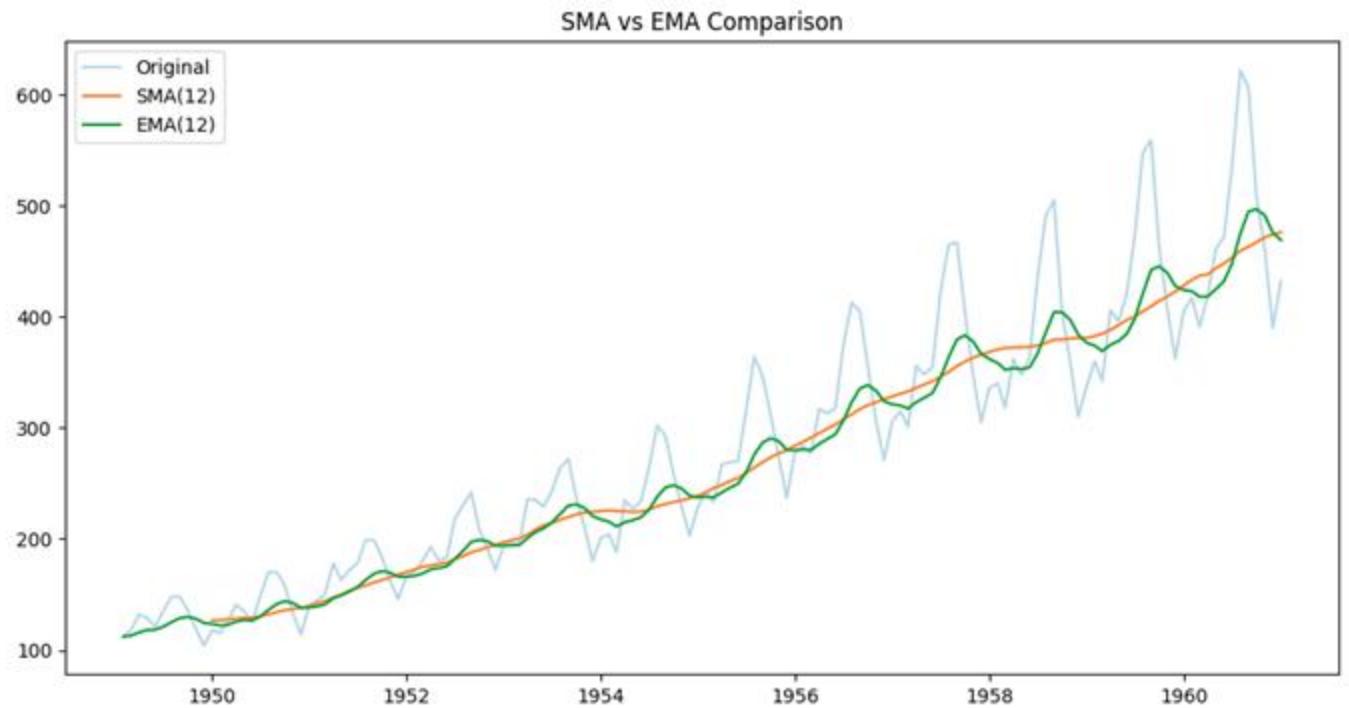
# Technical Indicators and EDA

## From SMA to EMA

```
# SMA
sma_12 = ts.rolling(window=12).mean()

# EMA
ema_12 = ts.ewm(span=12, adjust=False).mean()

# Comparison plot
plt.figure(figsize=(12, 6))
plt.plot(ts.index, ts.values, label='Original',
alpha=0.3)
plt.plot(sma_12.index, sma_12.values, label='SMA(12)')
plt.plot(ema_12.index, ema_12.values, label='EMA(12)')
plt.legend()
plt.title('SMA vs EMA Comparison')
plt.show()
```



# Technical Indicators and EDA

## Moving Average Convergence Divergence (MACD) Indicator

- Components:
  - Fast EMA (typically 12 periods)
  - Slow EMA (typically 26 periods)
  - MACD line = Fast EMA - Slow EMA
  - Signal line = EMA(9) of MACD line
- Interpretation:
  - Momentum indicator
  - MACD > Signal: Bullish momentum
  - MACD < Signal: Bearish momentum

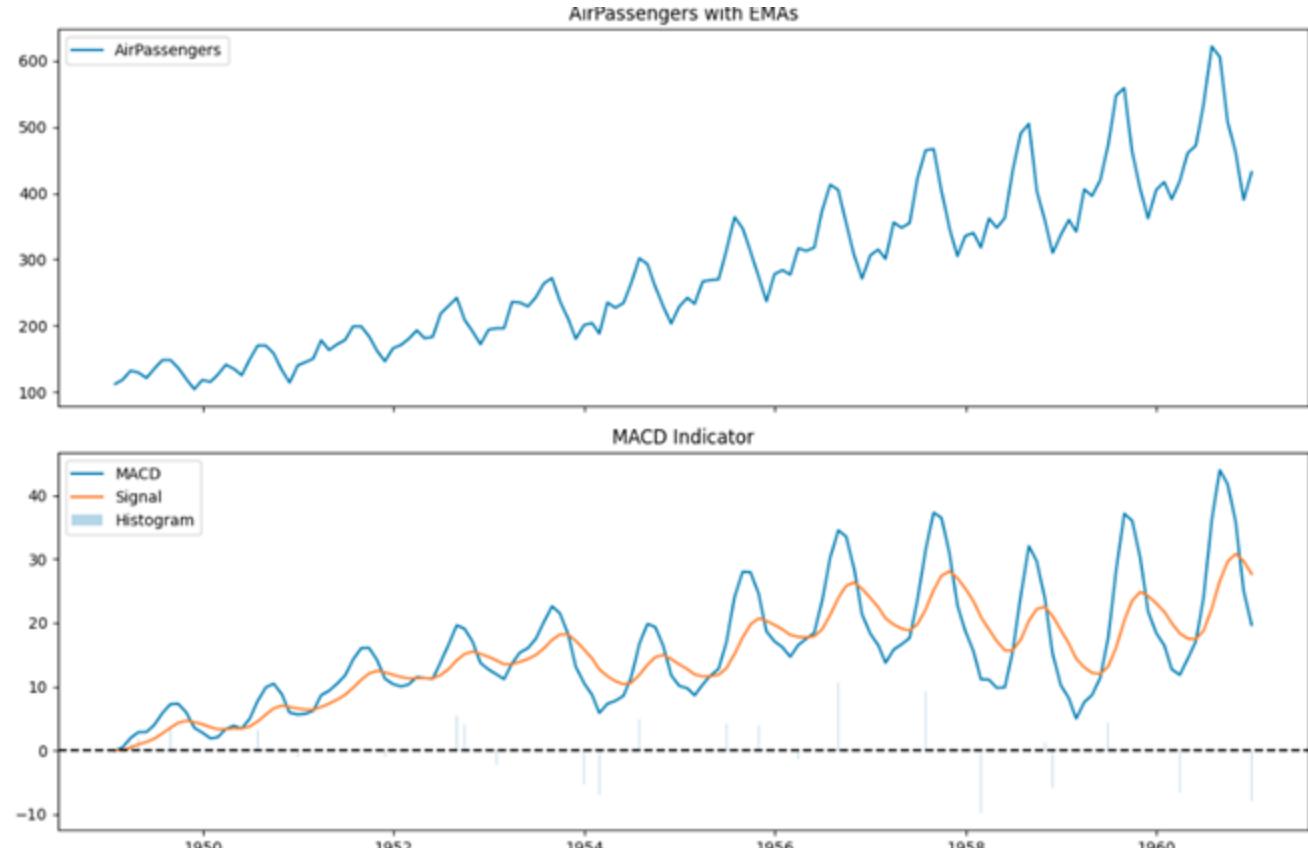
# Technical Indicators and EDA

## MACD Indicator

```
# Calculate MACD
ema_fast = ts.ewm(span=12, adjust=False).mean()
ema_slow = ts.ewm(span=26, adjust=False).mean()
macd_line = ema_fast - ema_slow
signal_line = macd_line.ewm(span=9, adjust=False).mean()
histogram = macd_line - signal_line

# Visualization
fig, axes = plt.subplots(2, 1, figsize=(12, 8), sharex=True)
axes[0].plot(ts.index, ts.values, label='AirPassengers')
axes[0].set_title('AirPassengers with EMAs')
axes[0].legend()

axes[1].plot(macd_line.index, macd_line.values, label='MACD')
axes[1].plot(signal_line.index, signal_line.values, label='Signal')
axes[1].bar(histogram.index, histogram.values, alpha=0.3, label='Histogram')
axes[1].axhline(y=0, color='black', linestyle='--')
axes[1].set_title('MACD Indicator')
axes[1].legend()
plt.tight_layout()
plt.show()
```



# Technical Indicators and EDA

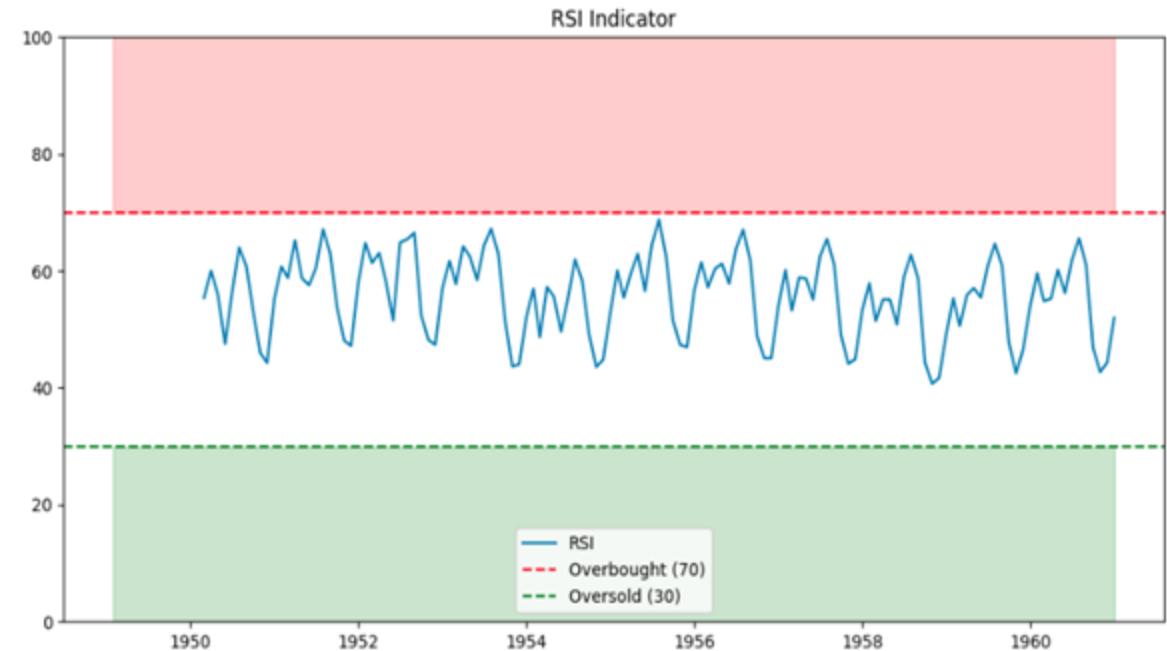
## Relative Strength Index (RSI) Indicator

- Relative Strength Index (RSI):
  - Momentum oscillator measuring speed and magnitude of changes
  - Formula:  $RSI = 100 \cdot (100 / (1 + RS))$   
Where  $RS = \text{Average Gain} / \text{Average Loss}$  (over n periods)
- Calculation:
  - Gains and losses over a window (typically 14 periods)
  - Separates positive and negative price changes
- Interpretation:
  - $RSI > 70$ : Overbought (potential reversal)
  - $RSI < 30$ : Oversold (potential reversal)
  - $RSI = 50$ : Neutral
- NonFinancial Illustration Purpose:
  - Demonstrates momentum concepts applicable to any time series
  - Useful for identifying extreme values in any domain

# Technical Indicators and EDA

## RSI Indicator (Conceptual)

```
def calculate_rsi(data, window=14):  
    delta = data.diff()  
    gain = (delta.where(delta > 0, 0)).rolling(window=window).mean()  
    loss = (-delta.where(delta < 0, 0)).rolling(window=window).mean()  
    rs = gain / loss  
    rsi = 100 - (100 / (1 + rs))  
    return rsi  
  
rsi = calculate_rsi(ts, window=14)  
  
# Visualization  
plt.figure(figsize=(12, 6))  
plt.plot(rsi.index, rsi.values, label='RSI')  
plt.axhline(y=70, color='r', linestyle='--', label='Overbought (70)')  
plt.axhline(y=30, color='g', linestyle='--', label='Oversold (30)')  
plt.fill_between(rsi.index, 70, 100, alpha=0.2, color='red')  
plt.fill_between(rsi.index, 0, 30, alpha=0.2, color='green')  
plt.ylim(0, 100)  
plt.title('RSI Indicator')  
plt.legend()  
plt.show()
```



# Technical Indicators and EDA

## Exploratory Data Analysis (EDA)

- Line plot of the series
  - Seasonal patterns identification
  - Rolling mean and variance analysis
  - Trend–seasonality decomposition
- 
- Excercise: proceed EDA with AirPassengers dataset

# Thank you!